ESL Exercises*

Chapter 2

Ex. 2.1 This result holds for standard Euclidean norm. For example we could consider a metric space $(\mathbb{R}^3, \|\cdot\|)$ where $\|a_0e_0 + a_1e_1 + a_2e_2\| = 3a_0^2 + a_1^2 + a_2^2$ and $\hat{y} = (0.4, 0.3, 0.3)$.

- Task 1: largest element of $\hat{y}$ is the first element.
- Task 2: $\|\hat{y} - t_0\| = 1.26$, $\|\hat{y} - t_{1,2}\| = 1.06$. $\arg\min \|\hat{y} - t_i\| = 1, 2 \neq 0$.

For standard Euclidean norm, we have,

$$\|t_k - \hat{y}\| = \sum_{i \neq k} \hat{y}_i^2 + (1 - \hat{y}_k)^2 = 1 + TSS(\hat{y}) - 2y_k.$$

Thus we have,

$$\begin{aligned} & \arg\min \|t_k - \hat{y}\| \\ = \ & \arg\min(1 + TSS(\hat{y}) - 2y_k) \\ = \ & \arg\min(-2y_k) \\ = \ & \arg\max(\hat{y}). \end{aligned}$$

Ex. 2.2 By the caption of Figure 2.5, we know $\mathbb{P}(x|C_k)$ and $\mathbb{P}(C_k)$. The Bayes decision boundary of the 0 - 1 loss is determined by, $\mathbb{P}(C_0|x) = \mathbb{P}(C_1|x)$ or equivalently

$$\mathbb{P}(x|C_0)\mathbb{P}(C_0) = \mathbb{P}(C_1|x)\mathbb{P}(C_1).$$

Ex. 2.3 $X = \min_{dist}\{X_1, ..., X_N\} \in \mathbb{R}^1$, $1 - F_X(x) = \mathbb{P}(X > x) = (1 - x^p)^N$. Median of X is simply the x_0 where $F_X(x_0) = \frac{1}{2}$.

Ex. 2.4 For samples follows multivariate Gaussian, $r^2 = \sum_{i=1}^p x_i^2 \sim \chi_p^2$. Fix a direction a, projection $x \cdot a = \sum_{i=1}^p x_i a_i \sim \mathcal{N}(0, \sum_{i=1}^p a_i^2) = \mathcal{N}(0, 1)$. (one can easily prove this using characteristic functions) Thus after projection, the mean is closer to zero in the $\mathbb{L}^2$ sense.

*https://github.com/xincui-math

Ex. 2.5

$$\begin{aligned} EPE(x_0) &= \mathbb{E}_{\mathcal{T},x_0}(y_0 - \hat{y_0})^2 \\ &= \mathbb{E}_{\mathcal{T},x_0}(y_0 - \mathbb{E}_{\mathcal{T},x_0}y_0)^2 + \mathbb{E}_{\mathcal{T},x_0}(\mathbb{E}_{\mathcal{T},x_0}y_0 - \hat{y_0})^2 \\ &= \mathbb{E}_{\mathcal{T},x_0}\epsilon^2 + \mathbb{E}_{\mathcal{T},x_0}(x_0\beta - x_0\hat{\beta})^2 \\ &= \mathbb{E}\epsilon^2 + x_0\mathbb{E}_{\mathcal{T}}(\beta - \hat{\beta})^2 \\ &= \mathbb{E}\epsilon^2 + x_0\mathbb{E}_{\mathcal{T}}\left(\beta - (X^TX)^{-1}X^T(X\beta + \epsilon)\right)^2 \\ &= \mathbb{E}\epsilon^2 + x_0\mathrm{Var}\left((X^TX)^{-1}X^T\epsilon\right) \\ &= \sigma^2 + x_0(X^TX)^{-1}. \end{aligned}$$

Ex. 2.6 Decompose sample space $\Omega = \bigoplus_{x_i}\Omega_i = \bigoplus_i\{(x_i, y_{ij})\}$.

$$\begin{aligned} SSR(\Omega_i) &= \sum_j [y_{ij} - f_\theta(x_i)]^2 \\ &= \sum_j y_{ij}^2 - 2\sum_j y_{ij}f_\theta(x_i) + n_i f_\theta(x_i)^2 \\ &= n_i\left(f_\theta(x_i) - \frac{1}{n_i}\sum_j y_{ij}\right)^2 + \phi(y_{ij}). \end{aligned}$$

Hence the problem reduces to weighted least square weights n_i.

Ex. 2.7 (a) Representations

Linear regression:

$$\begin{aligned} \hat{f}(x_0) &= x_0\hat{\beta} \\ &= x_0(X^TX)^{-1}X^Ty \\ &= \sum_i [x_0(X^TX)^{-1}X^T]_i y_i. \end{aligned}$$

K-nearest neighbourhood:

$$\hat{f}(x_0) = \sum_i \frac{1}{k} I_{i\in argmin_k \overrightarrow{d}(x_0,\mathcal{X})} y_i.$$

(b) $\mathbb{E}_{\mathcal{Y}|\mathcal{X}}(MSE)$

$$\begin{aligned}
& \mathbb{E}_{\mathcal{Y}|\mathcal{X}}\left[f(x_0)-\hat{f}(x_0)\right]^2 \\
= \; & \mathbb{E}_{\mathcal{Y}|\mathcal{X}}\left[f(x_0)-\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)+\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)-\hat{f}(x_0)\right]^2 \\
= \; & \mathbb{E}_{\mathcal{Y}|\mathcal{X}}\left[f(x_0)-\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)\right]^2+\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\left[\hat{f}(x_0)-\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)\right]^2 \\
& -2\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\left[f(x_0)-\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)\right]\left[\hat{f}(x_0)-\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)\right] \\
= \; & \left[f(x_0)-\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)\right]^2+\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\left[\hat{f}(x_0)-\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\hat{f}(x_0)\right]^2 \\
= \; & \left[f(x_0)-\sum_{i=1}^{N} l_i(x_0,\mathcal{X})f(x_i)\right]^2+\mathbb{E}_{\mathcal{Y}|\mathcal{X}}\left[\sum_{i=1}^{N} l_i(x_0,\mathcal{X})\epsilon_i\right]^2 \\
= \; & \left[f(x_0)-\sum_{i=1}^{N} l_i(x_0,\mathcal{X})f(x_i)\right]^2+\sum_{i=1}^{N} l_i^2(x_0,\mathcal{X})\sigma^2.
\end{aligned}$$

(c) $\mathbb{E}_{\mathcal{Y},\mathcal{X}}(MSE)$

Notice that $\hat{f}(x_0)$ is $(\mathcal{Y},\mathcal{X})$ measurable, $\mathbb{E}_{\mathcal{Y},\mathcal{X}}\hat{f}(x_0)=\hat{f}(x_0)$.

$$\begin{aligned}
& \mathbb{E}_{\mathcal{Y},\mathcal{X}}\left[f(x_0)-\hat{f}(x_0)\right]^2 \\
= \; & \left[f(x_0)-\mathbb{E}_{\mathcal{Y},\mathcal{X}}\hat{f}(x_0)\right]^2+\mathbb{E}_{\mathcal{Y},\mathcal{X}}\left[\hat{f}(x_0)-\mathbb{E}_{\mathcal{Y},\mathcal{X}}\hat{f}(x_0)\right]^2 \\
= \; & \left[f(x_0)-\hat{f}(x_0)\right]^2.
\end{aligned}$$

Ex. 2.8 TODO: code up.

Ex. 2.9 Write train set as (X_0,y_0), test set as (X_1,y_1), projection matrix as P_i. Rewrite $\hat{\beta}=(X_0^TX_0)^{-1}X_0^T\epsilon_0+\beta$. On train set, we have the following,

$$\begin{aligned}
& \mathbb{E}_0\left[(y_0-X_0\hat{\beta})^T(y_0-X_0\hat{\beta})\right] \\
= \; & \mathbb{E}_0\left[(\epsilon_0-P_0\epsilon)^T(\epsilon_0-P_0\epsilon_0)\right] \\
= \; & (N-k)\sigma^2.
\end{aligned}$$

On test set, we have,

$$\begin{aligned}
& \mathbb{E}_1\left[(y_1-X_1\hat{\beta})^T(y_1-X_1\hat{\beta})\right] \\
= \; & \mathbb{E}_1\left[(\epsilon_1-X_1(X_0^TX_0)^{-1}X_0^T\epsilon_0)^T(\epsilon_1-X_1(X_0^TX_0)^{-1}X_0^T\epsilon_0)\right] \\
= \; & N\sigma^2+trace\left[X_0(X_0^TX_0)^{-1}X_1^TX_1(X_0^TX_0)^{-1}X_0^T\right]\sigma^2 \\
= \; & N\sigma^2+trace\left[X_1(X_0^TX_0)^{-1}X_1^T\right]\sigma^2 \\
\geq \; & N\sigma^2.
\end{aligned}$$

Chapter 3

Ex. 3.1 For simplicity, let's denote variable with tilde as skipping skipping column i. Here we explore bit more in this problem to clarify how to compute F-score.

- RSS_0: $(y - X\beta)^T(y - X\beta)$
- RSS_1: $(y - \tilde{X}\tilde{\beta})^T(y - \tilde{X}\tilde{\beta})$
- rss_1: $(y - X\beta + X_i\beta_i)^T(y - X\beta + X_i\beta_i)$

or say, RSS_1 uses refit $\tilde{\beta}$, rss_1 uses original β.

$$rss_1 - RSS_0 = X_i^T X_i \beta_i^2 = (X^T X)_{ii}\beta_i^2.$$

Denote

- $span\tilde{X} = span\{X_k | k \neq i\}$
- $spanX = span\{X_k\}$
- P_A, projection matrix to $spanA$.
- $\tilde{X}^\perp$ is the orthogonal complement of $span\{\tilde{X}\}$ inside $span\{X\}$.

$$RSS_1 - RSS_0 = y^T P_{\tilde{X}} y - y^T P_X y = y^T P_{\tilde{X}^\perp} y.$$

$$F_i = \frac{(RSS_1 - RSS_0)/(p_1 - p_0)}{RSS_1/(N - p_1 - 1)} = \frac{y^T P_{\tilde{X}^\perp} y}{\hat{\sigma}^2} = \frac{x_i^T P_{\tilde{X}^\perp} x_i}{\hat{\sigma}^2}\hat{\beta i}^2 = \frac{\|X_i^\perp\|^2}{\hat{\sigma}^2}\hat{\beta i}^2.$$

$$z_i^2 = \frac{\hat{\beta}_i{}^2}{\hat{\sigma}^2 (X^T X)_{ii}^{-1}}.$$

Now let's compute $(X^T X)_{ij}^{-1}$. For arbitrary $\epsilon \sim \mathcal{N}(0, I_N)$,

$$\begin{aligned} & (X^T X)_{ij}^{-1} \\ = \; & COV\left((X^T X)^{-1} X^T \epsilon, (X^T X)^{-1} X^T \epsilon\right) \\ = \; & \mathbb{E}(\beta_{\epsilon,X,i}\beta_{\epsilon,X,j}) \\ = \; & \frac{X_i^\perp \cdot X_j^\perp}{\|X_i^\perp\|^2 \|X_j^\perp\|^2} \end{aligned}$$

Thus we have $z_i^2 = \frac{\hat{\beta}_i{}^2}{\hat{\sigma}^2 (X^T X)_{ii}^{-1}} = \frac{\hat{\beta}_i{}^2 \|X_i^\perp\|^2}{\hat{\sigma}^2} = F_i$. And $f_i = \frac{\hat{\beta}_i{}^2 (X^t X)_{ii}}{\hat{\sigma}^2}$

Ex. 3.2

- $\mathrm{Var}(a^T\beta) = a^T COV_\beta a = (X^T X)^{-1}\sigma^2$.
- $C_\beta = \{(\hat{\beta} - \beta)^T (X^T X)(\hat{\beta} - \beta) \leq \hat{\sigma}^2 \chi^2_{p+1}{}^{(1-\alpha)}\}$

The first estimation needs σ, the second one doesn't. (code TBD)

Ex. 3.3 For quantity $a^T\beta$, we have unbias estimator $\hat{\theta} = a^T\hat{\beta} = a^T(X^TX)^{-1}X^Ty$. Giving another unbiased estimation c^Ty, write $c^T = a^T(X^TX)^{-1}X^T + D$.

$$\mathbb{E}(\left[a^T(X^TX)^{-1}X^T + D\right]y) = a^T\beta + DX\beta.$$

Hence we have $DX = 0$.

$$\begin{aligned} & \mathrm{Var}(c^Ty) \\ = \ & \left[a^T(X^TX)^{-1}X^T + D\right]\left[a^T(X^TX)^{-1}X^T + D\right]^T\sigma^2 \\ = \ & [a^T(X^TX)^{-1}a + DD^T]\sigma^2 \\ \succeq \ & a^T(X^TX)^{-1}a\sigma^2 \\ = \ & \mathrm{Var}(a^T\hat{\beta}). \end{aligned}$$

Ex. 3.4 Give $X = (X_0, X_1, ..., X_{n-1})$ and y.

- $\tilde{\beta_0} = \frac{\mathrm{cov}(X_0,y)}{\mathrm{cov}(X_0,X_0)}$
- $\tilde{\beta_i} = \frac{\mathrm{cov}(z_i,y)}{\mathrm{cov}(z_i,z_i)}$ with $z_i = x_i - \sum_{j=0}^{i-1}\gamma_{ij}x_j$, $\gamma_{ij} = \frac{\mathrm{cov}(x_i,x_j)}{\mathrm{cov}(x_j,x_j)}$.

$$y = \sum_{i=0}^{n}\tilde{\beta_i}z_i = \sum_{i=1}^{n}\tilde{\beta_i}z_i = \sum_{i=1}^{n}\tilde{\beta_i}(x_i - \sum_{j=1}^{i-1}\gamma_{ij}x_j) = \sum_{i=1}^{n}(\tilde{\beta_i} - \sum_{j=i+1}^{n}\Gamma_{ji})x_i.$$

Ex. 3.5

- $\beta_0^c = \beta_0 + \sum_{j=0}^{p}\bar{x_j}\beta_j$
- $\beta_i^c = \beta_i$

Ex. 3.6 Assume prior $\beta \sim N(0, \tau I)$, data samples from $y \sim N(X\beta, \sigma^2 I)$. Posterior distribution has PDF proportional to:

$$p(\beta|D) \sim exp\left[-\frac{(y - X\beta)^T(y - X\beta)}{2\sigma^2} - \frac{\beta^T\beta}{2\tau^2}\right]$$

Q1: Equivalent to mode: $\lambda = \frac{\sigma^2}{\tau^2}$.
Q2: Equivalent to posterior mean:

$$-\frac{1}{2}(\beta - m_1)^T m_2^{-1}(\beta - m_1) = -\frac{(y - X\beta)^T(y - X\beta)}{2\sigma^2} - \frac{\beta^T\beta}{2\tau^2}$$

Solves the system, we have,

$$\begin{cases} m_1 = (\frac{X^TX}{\sigma^2} + \frac{I}{\tau^2}) \\ m_2 = (X^TX + \frac{\sigma^2}{\tau^2}X^Ty). \end{cases}$$

Ex. 3.7 Direct consequence of $p(\beta|D) \sim exp\left[-\frac{(y-X\beta)^T(y-X\beta)}{2\sigma^2} - \frac{\beta^T\beta}{2\tau^2}\right]$.

Ex. 3.8 For special matrix, $X = (e, x_1, x_2, x_3, ..., x_p)$, centered matrix is given by

$$\widetilde{X} = (x_1 - \overline{x_1}, ..., x_p - \overline{x_p}).$$

Gram-Schmidt processes gives the following. $x_i = \sum_{d=1}^{i-1} q_d r_{di} + \frac{\langle e, x_i \rangle}{|e|} \frac{e}{|e|}$.

$$Q_{lower} R_{lower} = \widetilde{X} = U \Sigma V^*.$$

Hence column span of Q_{lower} is the same as column span of U.

Denote $\tilde{R} = \Sigma^{-1} R = V^*$. Consider r_i be i-th column vector of $\tilde{R}$. Then use induction not hard to see $\tilde{R}$ is diagonal. This implies ΣV is diagonal of 1/-1.

Ex. 3.9 Denote $z_i = x_i - \sum_{k=1}^{r} q_k (x_i, q_k)$, variance explained increment has norm,

$$\|\hat{\beta}_i z_i\| = |\langle y, x_i - \sum_{k=1}^{r} q_k (x_i, q_k) \rangle|.$$

For new set of feature it is equivalent (and faster) to pick the following,

$$\text{argmax}_i \|(y^T X - y^T Q Q^T X)_i\|.$$

Ex. 3.10 Exercise 3.1 shows F statistics for dropping i-th variable corresponds to z_i^2. Hence we just need to drop the variable with lowest $|z|$.

Ex. 3.11

$$\begin{aligned}
& \left(\frac{dtr[(Y - XB)^T (Y - XB)]}{dB} \right)_{ij} \\
= & \frac{d}{db_{ij}} \sum_{p,q,v} (y_{pq} - X_{pv} B_{vq})^2 \\
= & -\sum_{p,q,v} 2(y_{pq} - X_{pv} B_{vq}) X_{ps} \delta_{sq}^{ij} \\
= & -\sum_{p,v} 2(y_{pj} - X_{pv} B_{vj}) X_{ps} \delta_{s}^{i} \\
= & -\sum_{p,v} 2(y_{pj} - X_{pv} B_{vj}) X_{pi} \\
= & -2(X^T Y - X^T X B)_{ij}.
\end{aligned}$$

Consider symmetric square root $\Sigma^{-\frac{1}{2}}$, the solution is

$$X^T Y \Sigma^{-\frac{1}{2}} - X^T X B \Sigma^{-\frac{1}{2}} = 0.$$

Ex. 3.12

$$\begin{pmatrix} y \\ 0 \end{pmatrix} = \begin{pmatrix} \widetilde{X} \\ \sqrt{\lambda} I \end{pmatrix} \beta + \epsilon.$$

$RSS(\beta) = (y - \tilde{X}\beta)^T (y - \tilde{X}\beta) + \lambda \|\beta\|^2$, same as Ridge regression.

Ex. 3.13 Now consider $z_i = Xv_i = \lambda_i u_i$.

- $\langle z_i, y\rangle = \lambda_i \langle u_i, y\rangle$.
- $\langle z_i, z_i\rangle = \lambda_i^2$.

$$\hat{\beta}^{pcr}(p) = \sum_{i=1}^{p} \frac{\langle z_i, y\rangle}{\langle z_i, z_i\rangle} v_i = VD^{-1}U^T y = \hat{\beta}^{ls}.$$

Ex. 3.14

- $z_1 = \sum_{j=1}^{p} \langle x_j, y\rangle x_j$.
- $x_j^1 = x_j^0 - \widehat{\phi_{1j}} \frac{z_1}{\langle z_1, z_1\rangle}$
- $\langle x_j^1, y\rangle = \langle x_j^0 - \widehat{\phi_{1j}} \frac{z_1}{\langle z_1, z_1\rangle}, y\rangle = \widehat{\phi_{1j}} - \widehat{\phi_{1j}} \frac{\langle z_1, y\rangle}{\langle z_1, z_1\rangle}$.

Using x_i are orthogonal, we have,

$$\langle z_1, y\rangle = \langle z_1, z_1\rangle = \sum_{j=1}^{p} \widehat{\phi_{1j}}^{\,2}.$$

This implies $\widehat{\phi_{2j}} = \langle x_j^1, y\rangle = 0$.

Ex. 3.15 (PLS)

$$\max_{\alpha} \operatorname{corr}^2(y, X\alpha) \operatorname{Var}(X\alpha)$$
$$\text{subject to } \|\alpha\| = 1, \alpha^T S \phi_l = 0.$$

The problem is equivalent to the following.

$$\max_{\alpha} \operatorname{cov}(y, X\alpha)$$
$$\text{subject to } \|\alpha\| = 1, \alpha^T S \phi_l = 0.$$

Decompose $\operatorname{span}(X) = X_{proj} \oplus X_{ortho}$ using inner product. Restrict α to $(X\alpha)_{proj} = 0$, Lagrange multiplier gives,

$$\begin{aligned} & L(x, \lambda) \\ = \; & \operatorname{cov}(y, X\alpha) - \lambda \alpha^T \alpha \\ = \; & \sum_i \left[\langle y, x_{iortho}\rangle \alpha_i - \lambda \alpha_i^2\right] + \langle y, (X\alpha)_{proj}\rangle. \end{aligned}$$

This implies $\alpha_i \sim \langle y, x_{iortho}\rangle$.

Ex. 3.16 Consider $y = \sum_i \alpha_i x_i + \epsilon$, where $\langle x_i, x_j \rangle = \delta_{ij}$.

Notice for any S, we always have estimated $\widehat{\beta_i^{(S)}} = \widehat{\beta}_i = \langle y, x_i \rangle$.

(1) Best M-subset

$$\begin{aligned} SSR(S) &= \left\| \sum_{i \in S} \widehat{\beta}_i x_i \right\|^2 \\ &= \sum_{i \in S} \widehat{\beta}_i{}^2. \end{aligned}$$

It is equivalent to pick the largest M $\widehat{\beta}_i$ in full regression.

(2) Ridge.

$$\begin{aligned} & \min\{\|y - X\beta\|^2 + \lambda \|\beta\|^2\} \\ = & \min_\beta \sum_i [(\lambda + 1)\beta_i^2 - 2\langle x_i, y \rangle \beta_i] \end{aligned}$$

Hence $\beta_i^{Ridge} = \frac{\widehat{\beta}_i}{\lambda + 1}$.

(3) Lasso.

$$\begin{aligned} & \min\{\|y - X\beta\|^2 + 2\lambda \|\beta\|\} \\ = & \min_\beta \sum_i \left[\beta_i^2 - 2[\hat{\beta}_i - \lambda \text{sign}(\beta_i)]\beta_i \right] \end{aligned}$$

(I) $\beta_i > 0$

- $\hat{\beta}_i \geq \lambda$: $\beta_i = \hat{\beta}_i - \lambda$
- $\hat{\beta}_i < \lambda$: $\beta_i = 0$

(II) $\beta_i < 0$

- $\hat{\beta}_i \geq -\lambda$: $\beta_i = 0$
- $\hat{\beta}_i < \lambda$: $\beta_i = \hat{\beta}_i + \lambda$

Rephrase the above analysis gives $\beta_j^{Lasso} = \text{sign}(\hat{\beta}_j)(\hat{\beta}_j - \lambda)_+$.

Ex. 3.17 Code TBD.

Ex. 3.18 Solving β is equivalent represent y_{proj} in coordinate X. The PLS solves a set of orthonormal basis iteratively for space span$\{X\}$, namely z_m. Under z_m,

$$\beta_{z_m} = \frac{\langle z_m, y \rangle}{\langle z_m, z_m \rangle_I}.$$

This formula matches the conjugate gradients algorithm.

Ex. 3.19 (1) L^2 norm as a decreasing function of λ in Ridge regression.

$$\begin{aligned}
& \frac{d}{d\lambda}\|\beta^{ridge}\|^2 \\
= & \frac{d}{d\lambda}[y^T X(X^TX+\lambda I)^{-1}(X^TX+\lambda I)^{-1}X^Ty] \\
= & -2\lambda[y^T X(X^TX+\lambda I)^{-1}(X^TX+\lambda I)^{-1}(X^TX+\lambda I)^{-1}X^Ty] \\
= & -2\lambda\beta_\lambda^T(X^TX+\lambda I)^{-1}\beta_\lambda \\
\leq & 0
\end{aligned}$$

(2) L^1 norm may not be a decreasing function of λ in Lasso regression.

Ex. 3.20 (CCR problem) Follow the exact same approach in [3]. $c = \Sigma_{YY}^{1/2}u$, $d = \Sigma_{XX}^{1/2}v$. Cauchy Schwarz gives,

$$u^TY^TXv \leq \frac{\|\Sigma_{XX}^{-1/2}\Sigma_{XY}\Sigma_{YY}^{-1/2}c\|}{\|c\|}.$$

Equality holds when $d = \lambda\Sigma_{XX}^{-1/2}\Sigma_{XY}\Sigma_{YY}^{-1/2}c$. Perform SVD decomposition on $\Sigma_{XX}^{-1/2}\Sigma_{XY}\Sigma_{YY}^{-1/2} = UDV^T$, and consider $\tilde{c} = V^Tc$.

$$\frac{\|\Sigma_{XX}^{-1/2}\Sigma_{XY}\Sigma_{YY}^{-1/2}c\|}{\|c\|} = \frac{\|\tilde{c}^T\Sigma\tilde{c}\|}{\|\tilde{c}\|}.$$

Optimization result $\tilde{c}$ gives identity matrix. Hence c gives V, right singular vectors, equality condition gives d are left singular vectors of $\Sigma_{XX}^{-1/2}\Sigma_{XY}\Sigma_{YY}^{-1/2}$. This gives same conclusion (after transpose).

$$\begin{aligned}
u_1 &= \Sigma_{YY}^{-1/2}u_1^* \\
v_1 &= \Sigma_{XX}^{-1/2}v_1^*
\end{aligned}$$

Ex. 3.21

$$\begin{aligned}
& tr\left[(y-XB)\Sigma^{-1}(y-XB)^T\right] \\
= & tr\left[(y^*-XB\Sigma^{-1/2})^T(y^*-XB\Sigma^{-1/2})\right] \\
= & tr\left[(Z-A)(Z-A)^T\right] + \text{const}(B)
\end{aligned}$$

where $Z = \Sigma_{YY}^{-1/2}\Sigma_{YX}\Sigma_{XX}^{-1/2}$, $A = \Sigma_{YY}^{-1/2}B^T\Sigma_{XX}^{1/2}$. Apply Theorem 3.7.4 in [4],

$$A_{opt} = \sum_{j=1}^{m} d_jU_jV_j^T.$$

where U, V are singular vectors of Z. Hence $B = \Sigma_{XX}^{-1}\Sigma_{YX}\sum_i u_{ccr,i}u_{ccr,i}^T$, where $u_{ccr,i} = \Sigma_{YY}^{-1/2}u_i$. U is orthgonal gives $U_m^T\Sigma_{YY}^{1/2}$ is generalized inverse of $\Sigma_{YY}^{-1/2}U_m$.

Ex. 3.22 Replace Σ_{YY} to $\Sigma_{residual}$.

Ex. 3.23 (a)

$$\begin{aligned} & \|\langle x_j, y - u(\alpha)\rangle\| \\ = \; & \|\langle x_j, y - \alpha X(X^TX)^{-1}X^Ty\rangle\| \\ = \; & \|(1-\alpha)(X^Ty)_j\| \\ = \; & N\lambda|1-\alpha|. \end{aligned}$$

(b) $(y - \alpha X\hat{\beta})^T(y - \alpha X\hat{\beta}) = N + \alpha(\alpha - 2)y^TX\hat{\beta}$. Let $\alpha = 1$, we have, $y^TX\hat{\beta} = N - RSS$. Hence $(y - \alpha X\hat{\beta})^T(y - \alpha X\hat{\beta}) = N(1-\alpha)^2 + \alpha(2-\alpha)RSS$.

$$\text{corr}(x_i, y - u(\alpha)) = \frac{\langle x_i, y - u(\alpha)\rangle}{\|x_i\|\|y - u(\alpha)\|} = \frac{N\lambda|1-\alpha|}{\sqrt{N}\sqrt{N(1-\alpha)^2 + \alpha(2-\alpha)RSS}}.$$

(c) $(X_{\mathcal{A}_k}^TX_{\mathcal{A}_k})^{-1}X_{\mathcal{A}_k}^Tr_k$ is OLS of r_k with $X_{\mathcal{A}_k}$. Notice that X has mean 0, std 1, $X_{\mathcal{A}_k}$ has same correlations with r_k. Hence (2) directly gives the result.

Ex. 3.24

$$\cos(x_i, X\hat{\beta}) = \frac{(X^TX\hat{\beta})_i}{\|x_i\|\|X\hat{\beta}\|} = \frac{(X_{\mathcal{A}_k}^Tr_k)_i}{\sqrt{N}\|X\hat{\beta}\|}.$$

$(X_{\mathcal{A}_k}^Tr_k)_i$ is constant among i from LAR algorithm.

Ex. 3.25 Equivalent to pick minimum α_i for tie covariance.

$$\min_{\alpha_i}\{N(1-\alpha)\lambda = |x_i^Tr_k - \alpha x_i^Tv_{\mathcal{A}_k}u|\}.$$

Ex. 3.26 Normalize $\widetilde{x_i} = (x_i - \sum_k \frac{\langle x_i, z_k\rangle}{\langle z_k, z_k\rangle}z_k)/\|x_i - \sum_k \frac{\langle x_i, z_k\rangle}{\langle z_k, z_k\rangle}z_k\|$. Notice that selecting direction $\widetilde{x_i}$ is equivalent to selecting x_i. SSR increment of picking $\widetilde{x_i}$ is $|\langle\widetilde{x_j}, \widetilde{r_k}\rangle|$.

Ex. 3.27 Objective $L(\beta) + \lambda\sum_j(\beta_j^+ + \beta_j^-)$, constrains $-\beta_j^\pm \le 0$. Hence dual function is, $L(\beta) + \lambda\sum_j(\beta_j^+ + \beta_j^-) - \sum_j\lambda_j^+\beta_j^+ - \sum_j\lambda_j^-\beta_j^-$. KKT condition.

$$\begin{aligned} \nabla L_j + \lambda - \lambda_j^+ &= 0 && \text{(Stationarity, } \beta_j^+\text{)} \\ -\nabla L_j + \lambda - \lambda_j^- &= 0 && \text{(Stationarity, } \beta_j^-\text{)} \\ \beta_j^\pm &\ge 0 && \text{(Primal feasibility)} \\ \lambda_j^\pm &\ge 0 && \text{(Dual feasibility)} \\ \lambda_j^\pm\beta_j^\pm &= 0 && \text{(Complementary slackness)} \end{aligned}$$

(b) $|\nabla L_j| = \frac{1}{2}(\lambda_j^+ - \lambda_j^-) \le \frac{1}{2}(\lambda_j^+ + \lambda_j^-) = \lambda$.
Case 1: $\lambda = 0$, inequality above implies $\nabla L_j = 0$.

Case 2: $\lambda > 0$, $\beta_j^+ > 0$, complementary slackness gives $\lambda_j^+ = 0$ hence $\lambda_j^- > 0$ which implies $\beta_j^- = 0$ again by complementary slackness. $\nabla L_j = -\lambda$.

Case 3: similar to case 2.

Combine the gradient, we have a expression of λ.

$$\nabla_j L = -x_j^T(y - X\beta_\lambda).$$

If active predictor is not changed $X^T X\beta_\lambda$ is an affine vector of λ, hence β_λ is affine.

Ex. 3.28 Consider $\beta = \beta_i^{(1,2)}$ solves duplicated Lasso optimization, we have $|\beta_i^{(1)}| + |\beta_i^{(2)}| \geq |\beta_i^{(1)} + \beta_i^{(2)}|$. This implies $\beta_i^{(1)} + \beta_i^{(2)}$ solves the original Lasso problem (3.51) with the exact same t. Give the original Lasso problem, any pair $\beta_i^{(1)} = \lambda\beta_i$, $\beta_i^{(2)} = (1-\lambda)\beta_i$ also solves the duplicated Lasso problem (for any λ in [0, 1]).

Ex. 3.29 Optimization of the duplicated ridge regression.

$$L(\beta) = (y - X\sum_i \beta_i)^T(y - X\sum_i \beta_i) + \lambda\sum_i \beta_i^T\beta_i.$$

$\nabla_{\beta_i} L(\beta) = -2X(y - X\sum_i \beta_i) + 2\lambda\beta_i = 0$. This implies $\beta_i = \beta_j = \beta$. Thus $\beta = (nX^TX + \lambda I)^{-1}X^Ty$.

Ex. 3.30

$$\begin{pmatrix} y \\ 0 \end{pmatrix} = \begin{pmatrix} \widetilde{X} \\ \sqrt{\alpha\lambda}I \end{pmatrix}\beta + \epsilon.$$

$\text{Lasso}_{\frac{\lambda(1-\alpha)}{2}}(\beta) = \frac{1}{2}(y - \tilde{X}\beta)^T(y - \tilde{X}\beta) + \frac{\alpha\lambda}{2}\|\beta\|^2 + \frac{\lambda(1-\alpha)}{2}\|\beta\|$.

References

[1] Hastie, T., Tibshirani, R.,, Friedman, J. (2001). The Elements of Statistical Learning. New York, NY, USA: Springer New York Inc..

[2] John L. Weatherwax, David Epstein A Solution Manual and Notes for: The Elements of Statistical Learning by Jerome Friedman, Trevor Hastie, and Robert Tibshirani, `https://waxworksmath.com/Authors/G_M/Hastie/WriteUp/Weatherwax_Epstein_Hastie_Solution_Manual.pdf`.

[3] Canonical correlation. Wikipedia: The Free Encyclopedia. Wikimedia Foundation, Inc. 30 Nov 2020. Web. 15 Feb. 2021, `https://en.wikipedia.org/wiki/Canonical_correlation`.

[4] Brillinger, David R. Time series: data analysis and theory. Society for Industrial and Applied Mathematics, 2001.